



At $\gamma = 0$: $\tanh \gamma = 0$, the gate is closed, $\mathbf{h}' = \mathbf{h}$, and the augmented model equals the base LM. Gradient reaches γ first ($\text{sech}^2 0 = 1$), then opens the visual path. Cf. LoRA's zero-initialized adapter.