



$\times N/8$ periods

$\vdots$



$$M = N[f \cdot 2wd + (1 - f) d_s d] b$$



Every layer contributes a decode-state term independent of n : $d_s d$ per recurrent block, $2wd$ for the one sliding-window attention layer per period. Stacking $N/8$ periods keeps total decode memory fixed in n , unlike full attention's per-layer cache $2nd$, which grows linearly.