




each $\text{head}_i = \text{Attn}(XW_i^Q, XW_i^K, XW_i^V) \in \mathbb{R}^{n \times d_v}$ — its *own* learned projections, not a slice of X
 h heads $\cdot$ concat $\cdot$ remix by W^O (mixes across heads) (example $d=6$, $h=3$, $d_k=d_v=2$)