


n visual features



Cross-attention $\text{CA}(\mathbf{Q}_{\text{lat}}, \mathbf{V})$: each soft token is a convex average of the projected values, and the output length is m for every n .